

# Learning Large-Scale Social Knowledge Graphs

Zhilin Yang, Jie Tang

Dept. Computer Science, Tsinghua University


# Large-scale social networks

- Facebook
  - 1.4 billion active users in Quarter 1, 2015
  - Tens of millions of posts per day
- AMiner
  - 39 million researchers
  - 79 million papers
- Large-Scale social networks are big information networks!

# Large-scale collective knowledge

- Freebase
  - 44 million entities
  - 2.4 billion facts
- YAGO2
  - 10 million entities
  - 120 million facts
- Wikipedia
  - 35 million entities
  - 2 million categories

# Bridge the gap



# Bridge the gap

- Social knowledge graph
- Why?
  - Better mine large volume of information
  - Better user understanding and recommendation
  - Better search

# What we've done

- Propose an algorithm GenVector to learn large-scale social knowledge graph
  - Weakly supervision based on unsupervised techniques
  - Multi-source Bayesian embedding model
- Online deployment
  - Online service on AMiner.org
  - Online AB-test

# Key features

- Large-scale
  - **38,049,189** researchers (AMiner)
  - **74,050,920** papers (AMiner)
  - **20,552,544,886** bytes corpus (Wikipedia full text)
  - **35,415,011** entities (Wikipedia)

# Key features

- Large-scale
- Fast
  - Implementation optimization for a **60 times** speedup
  - From 3 hours per iteration to 3 minutes



# Key features

- Large-scale
- Fast
- Accurate
  - Offline test: 4% to 15%+ better than state-of-the-arts
  - Online test: decrease the error rate by **67%**

# Key features

- Large-scale
- Fast
- Accurate
- Novel
  - Bridge the gap between social networks and collective knowledge
  - Bridge the gap between topic models and word/network embedding

# Key features

- Large-scale
- Fast
- Accurate
- Novel
- Real-world impact
  - Online deployment on AMiner
  - **183,876** visits ever since

# Key features

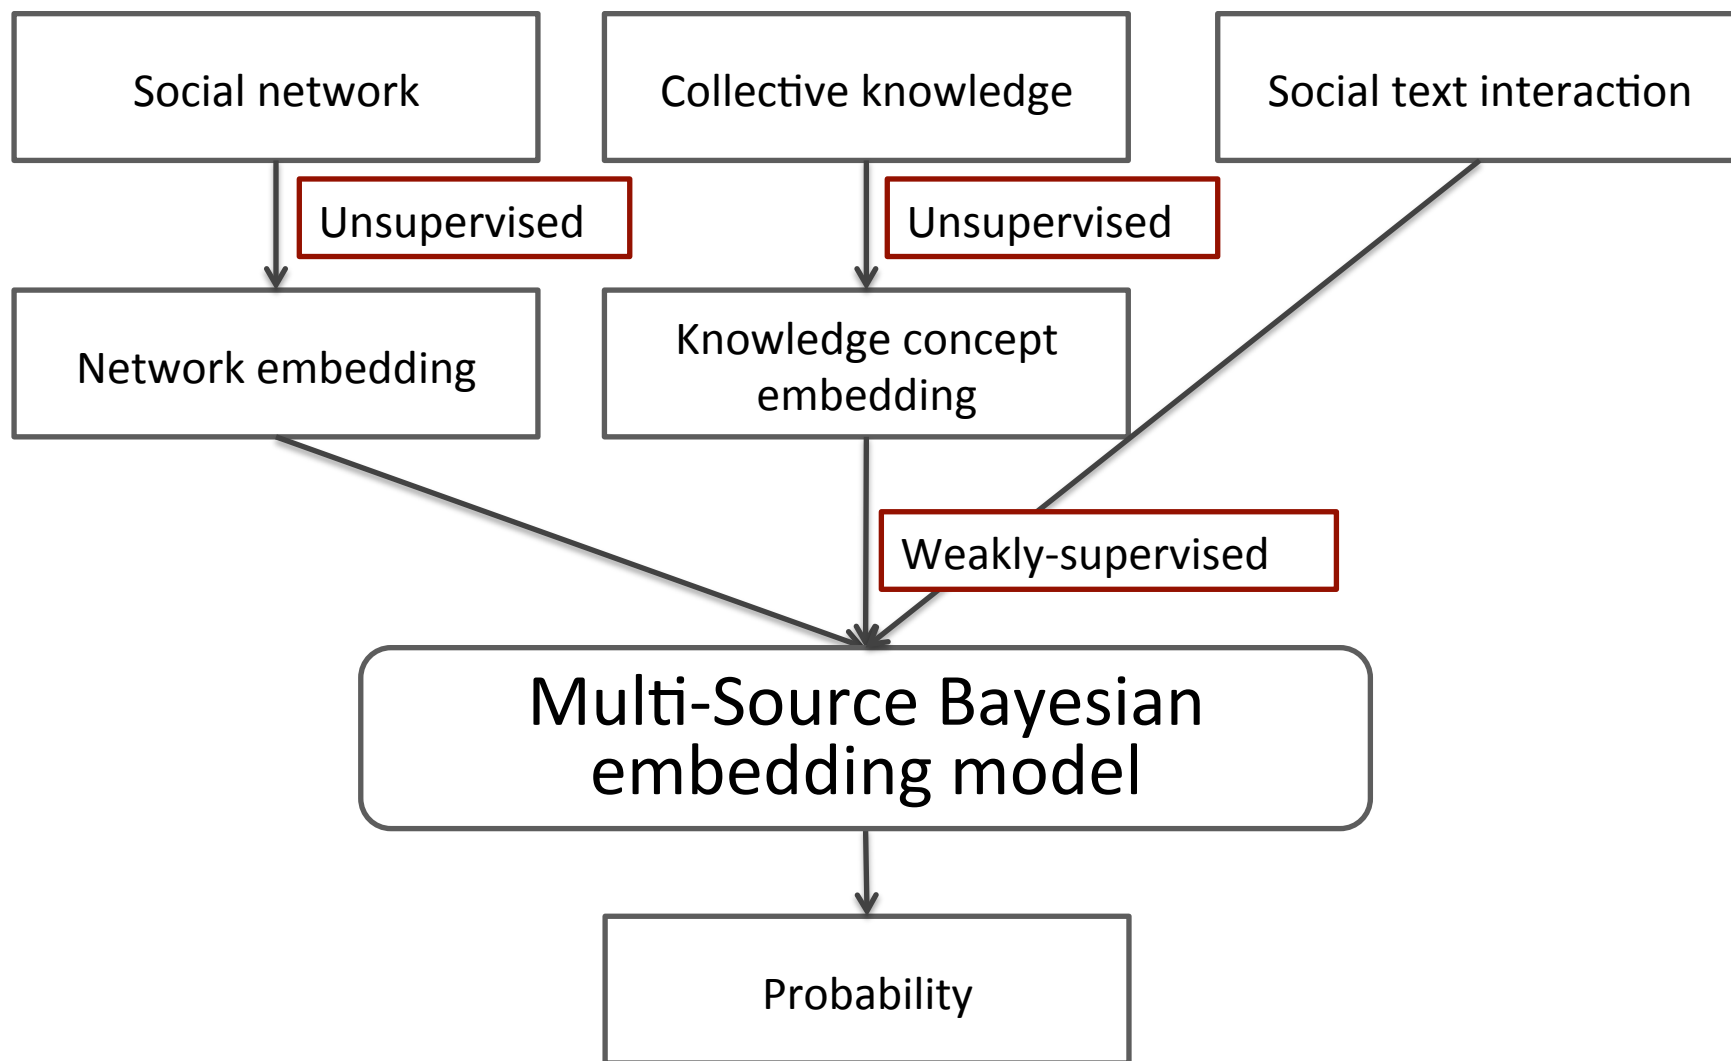
- Large-scale
- Fast
- Accurate
- Novel
- Real-world impact

**How did we make it?**

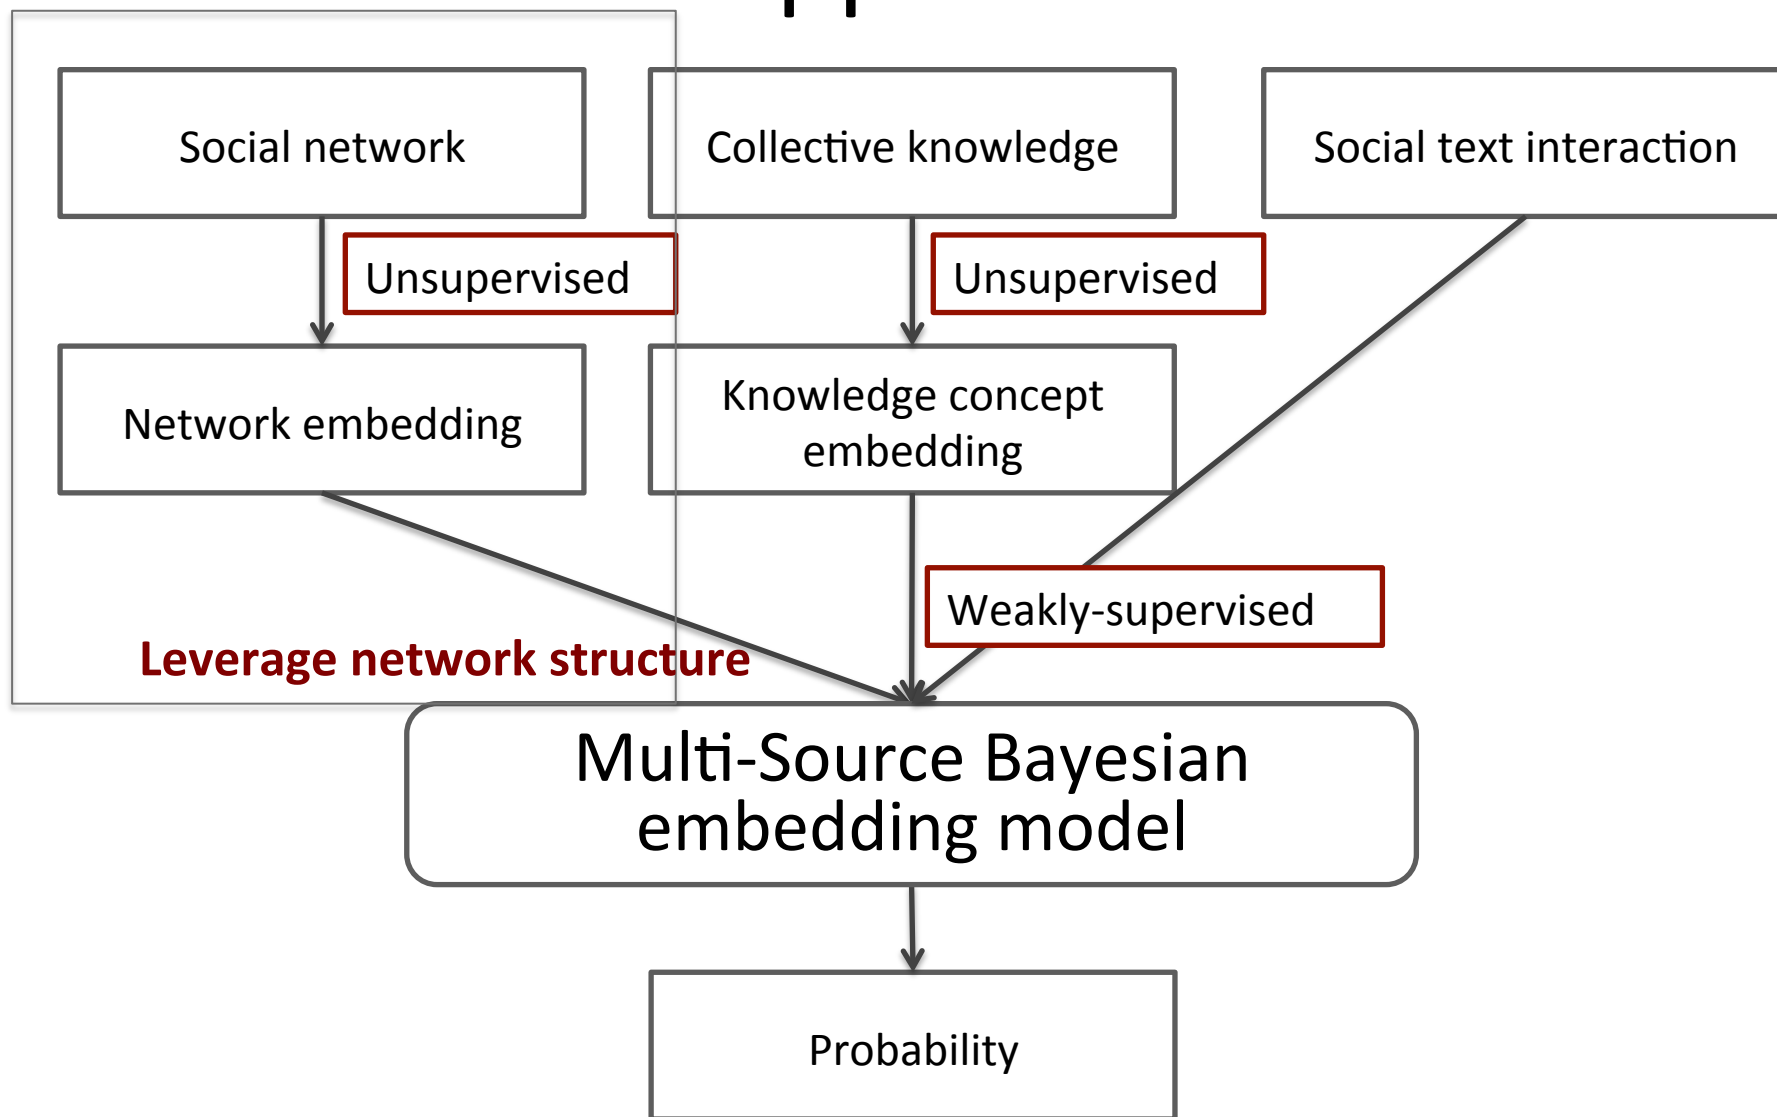
# Problem formulation

- Input
  - A social network
  - A collective knowledge source
  - Social text interaction
- Output
  - For each social network vertex, output related knowledge concepts as a ranked list

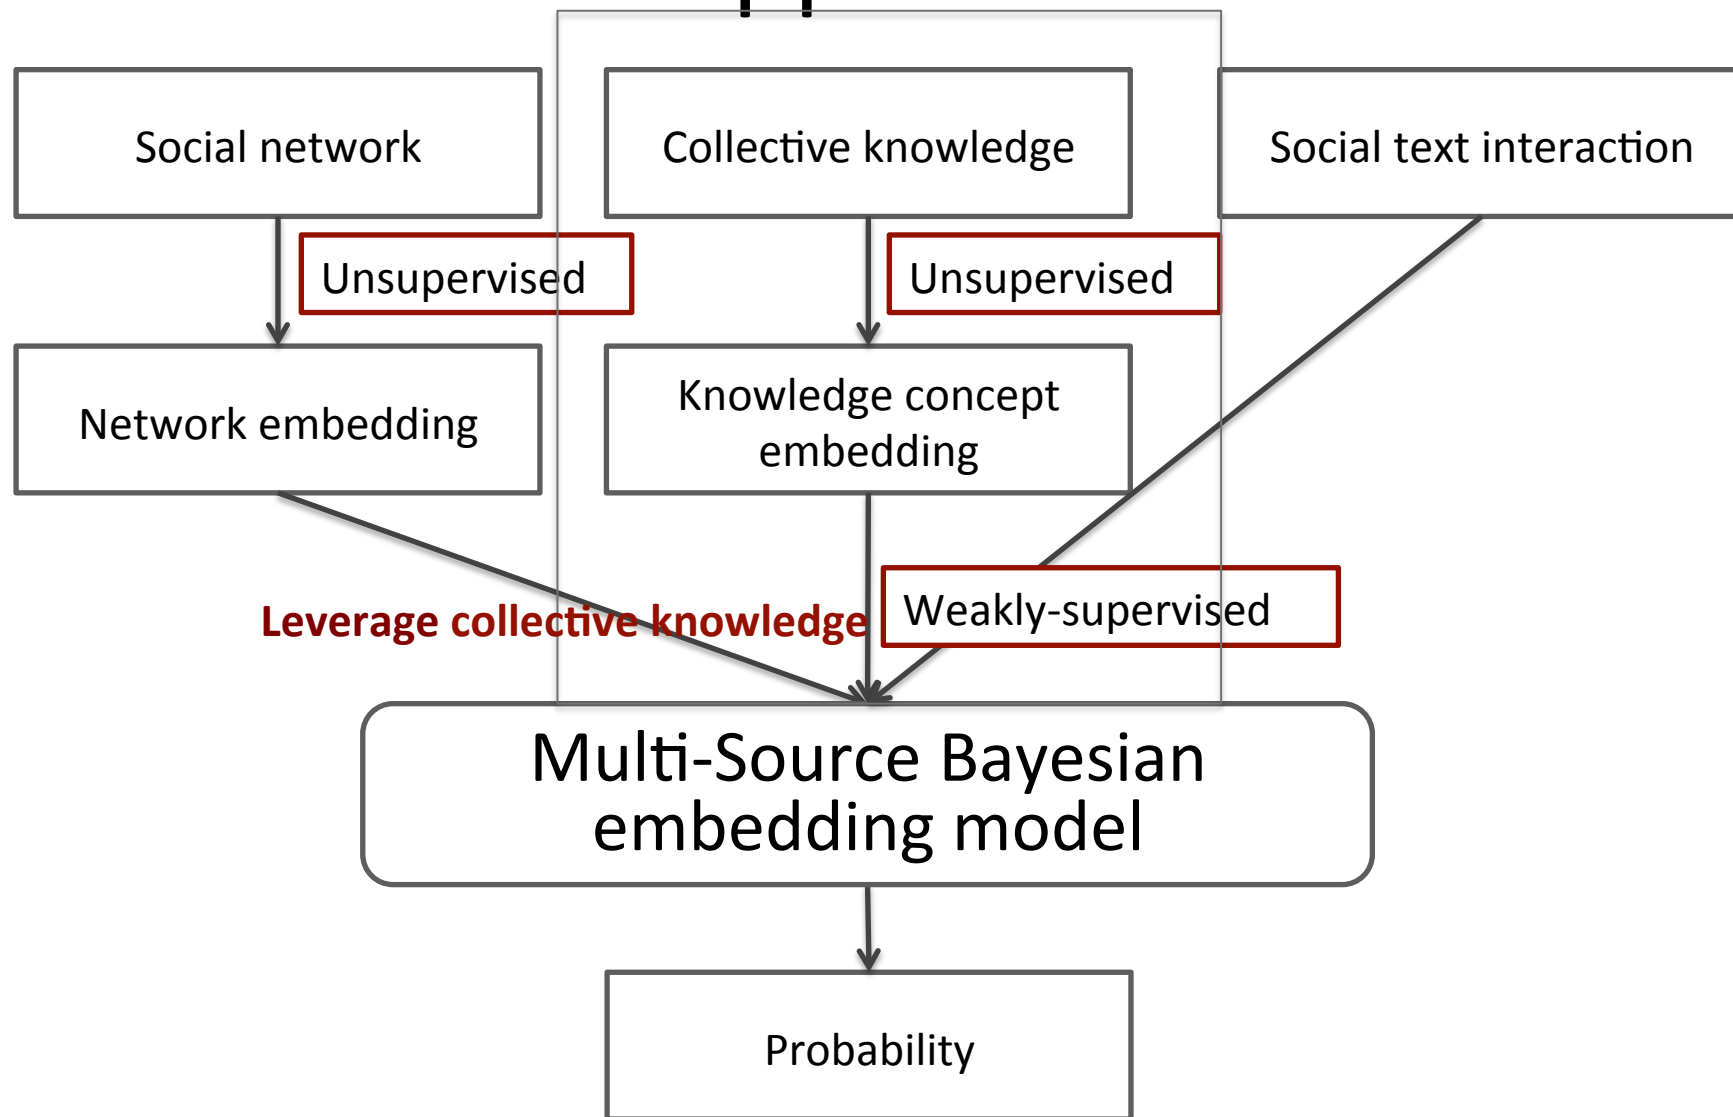
# Approach



# Approach

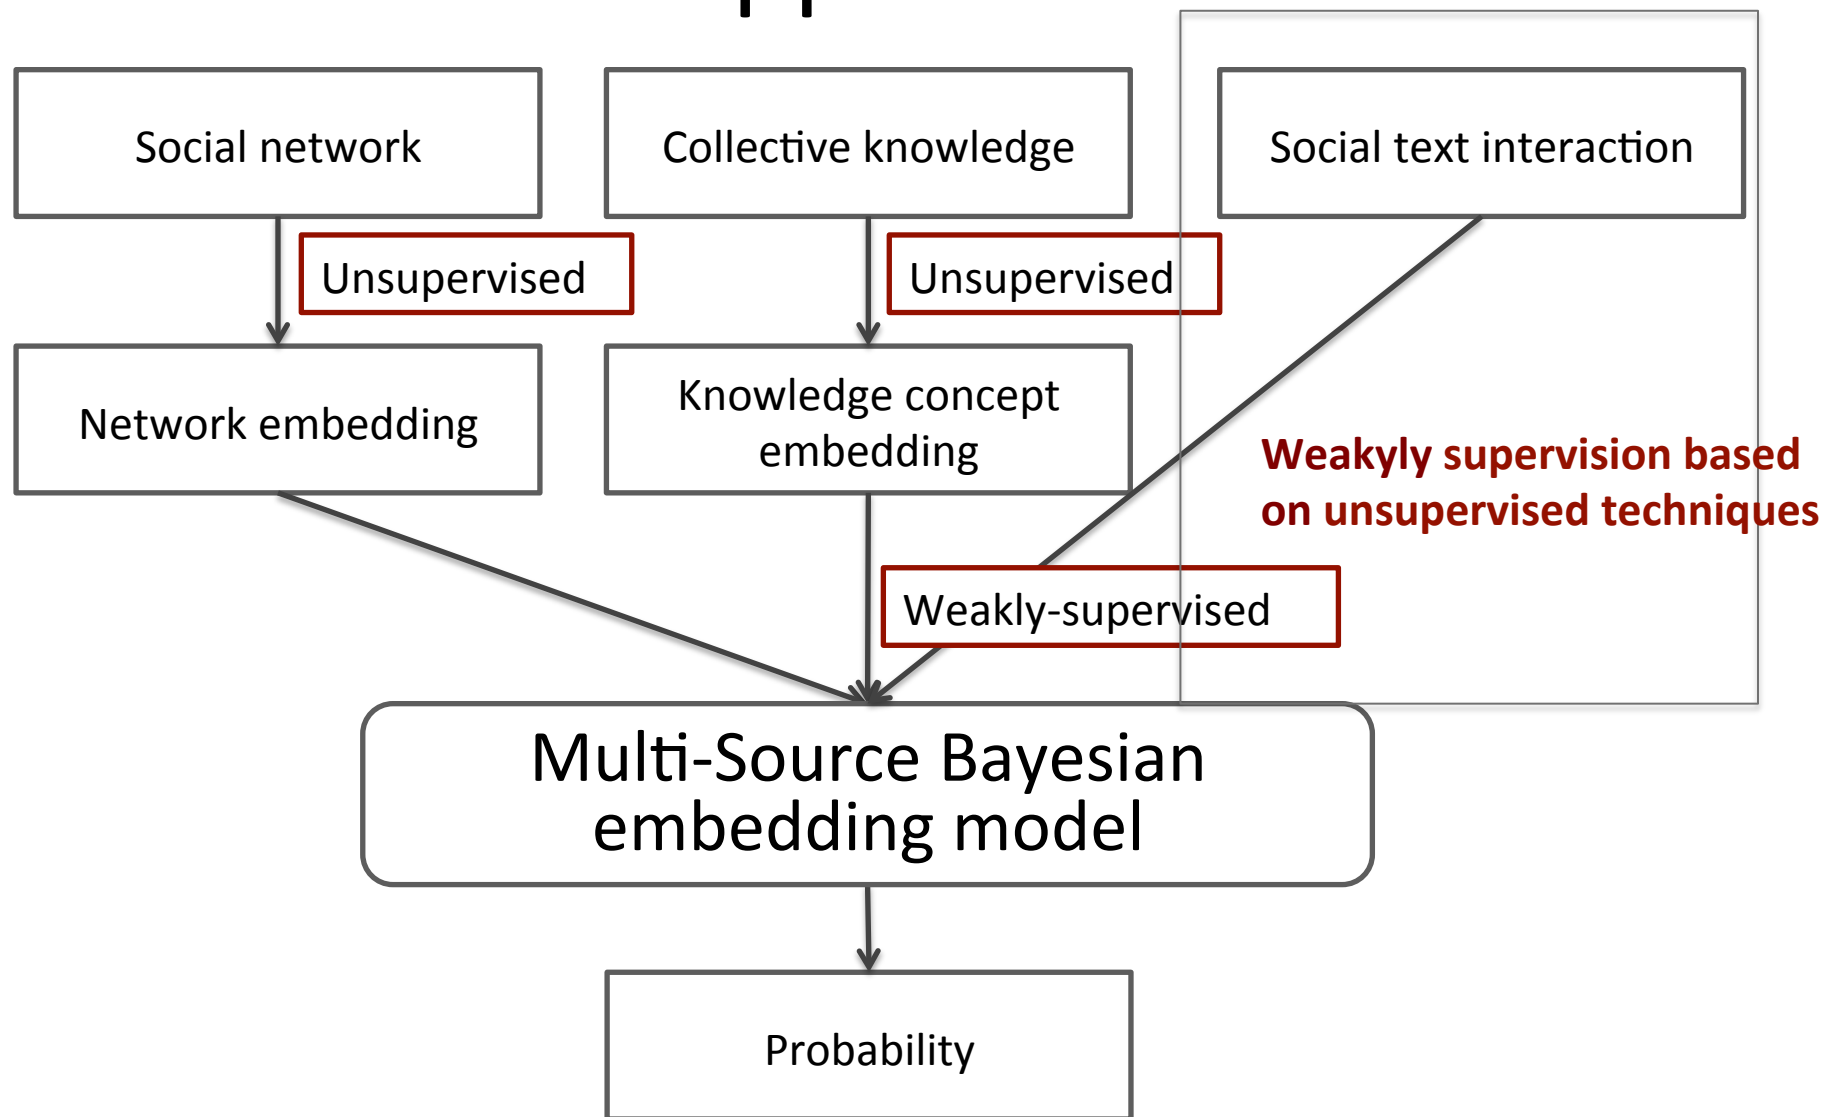


# Approach

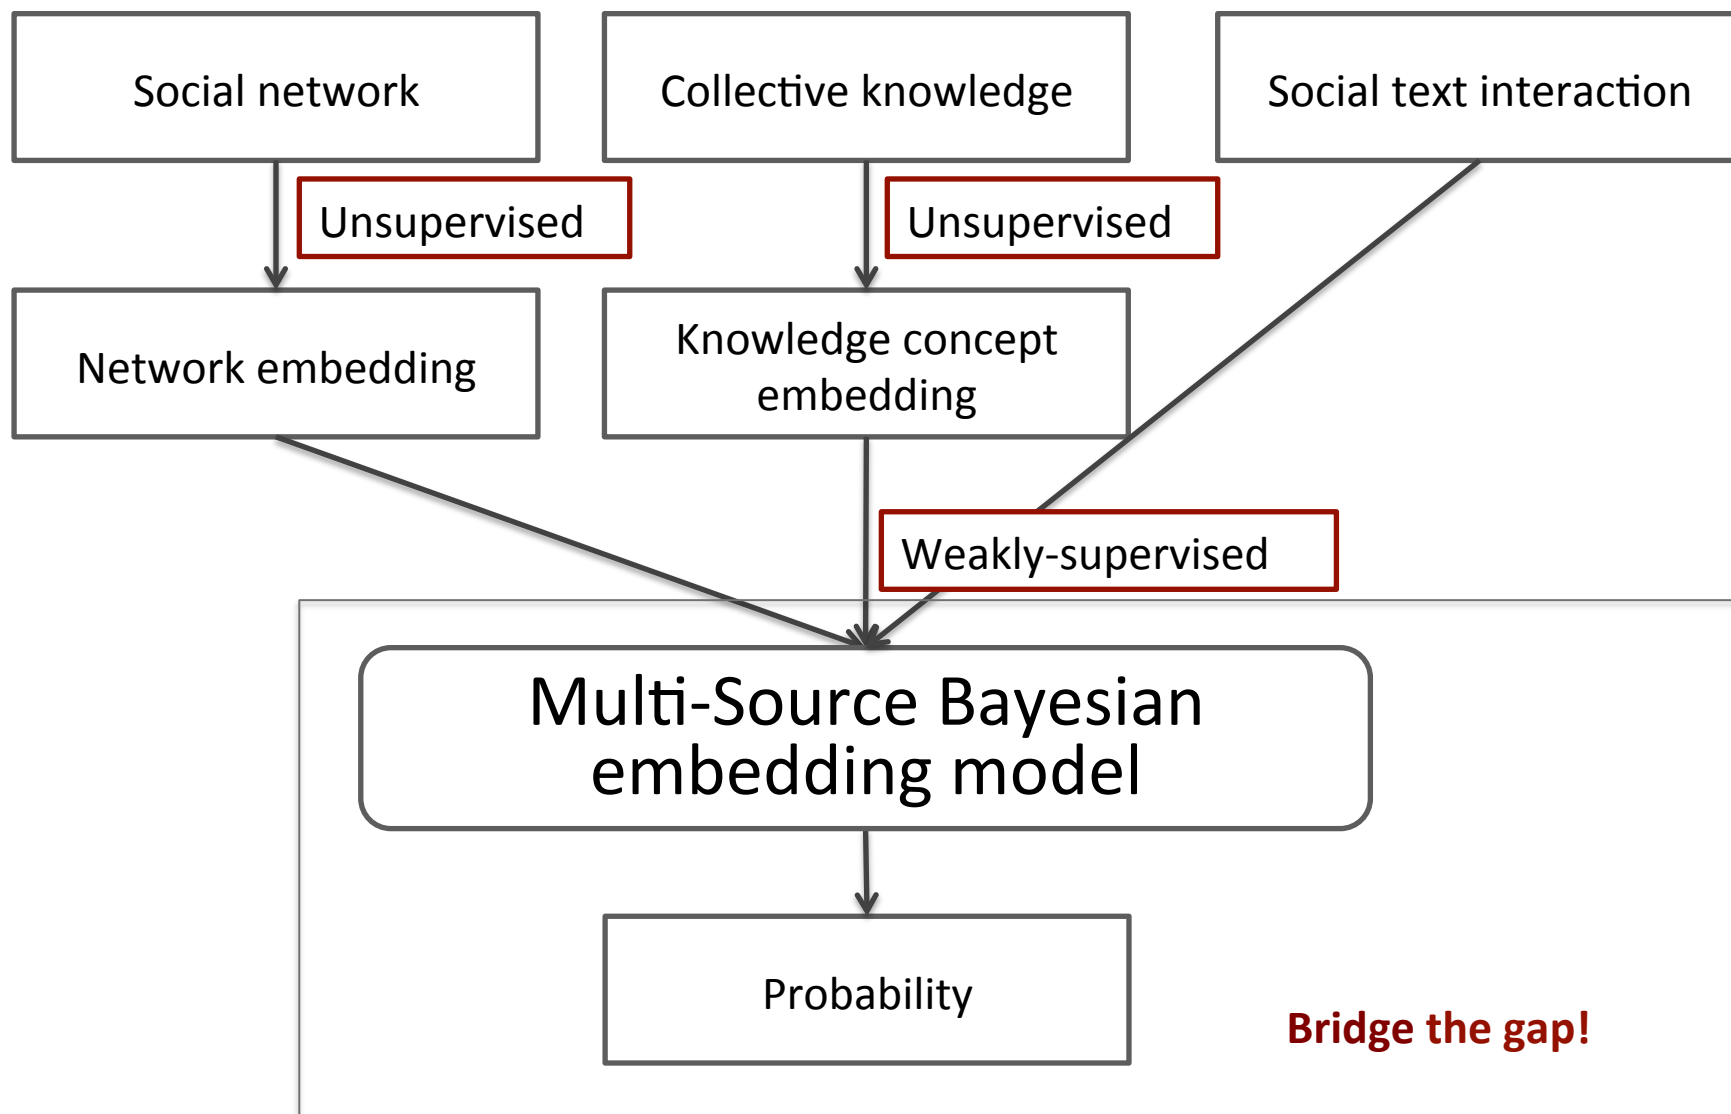





# Approach



# Approach

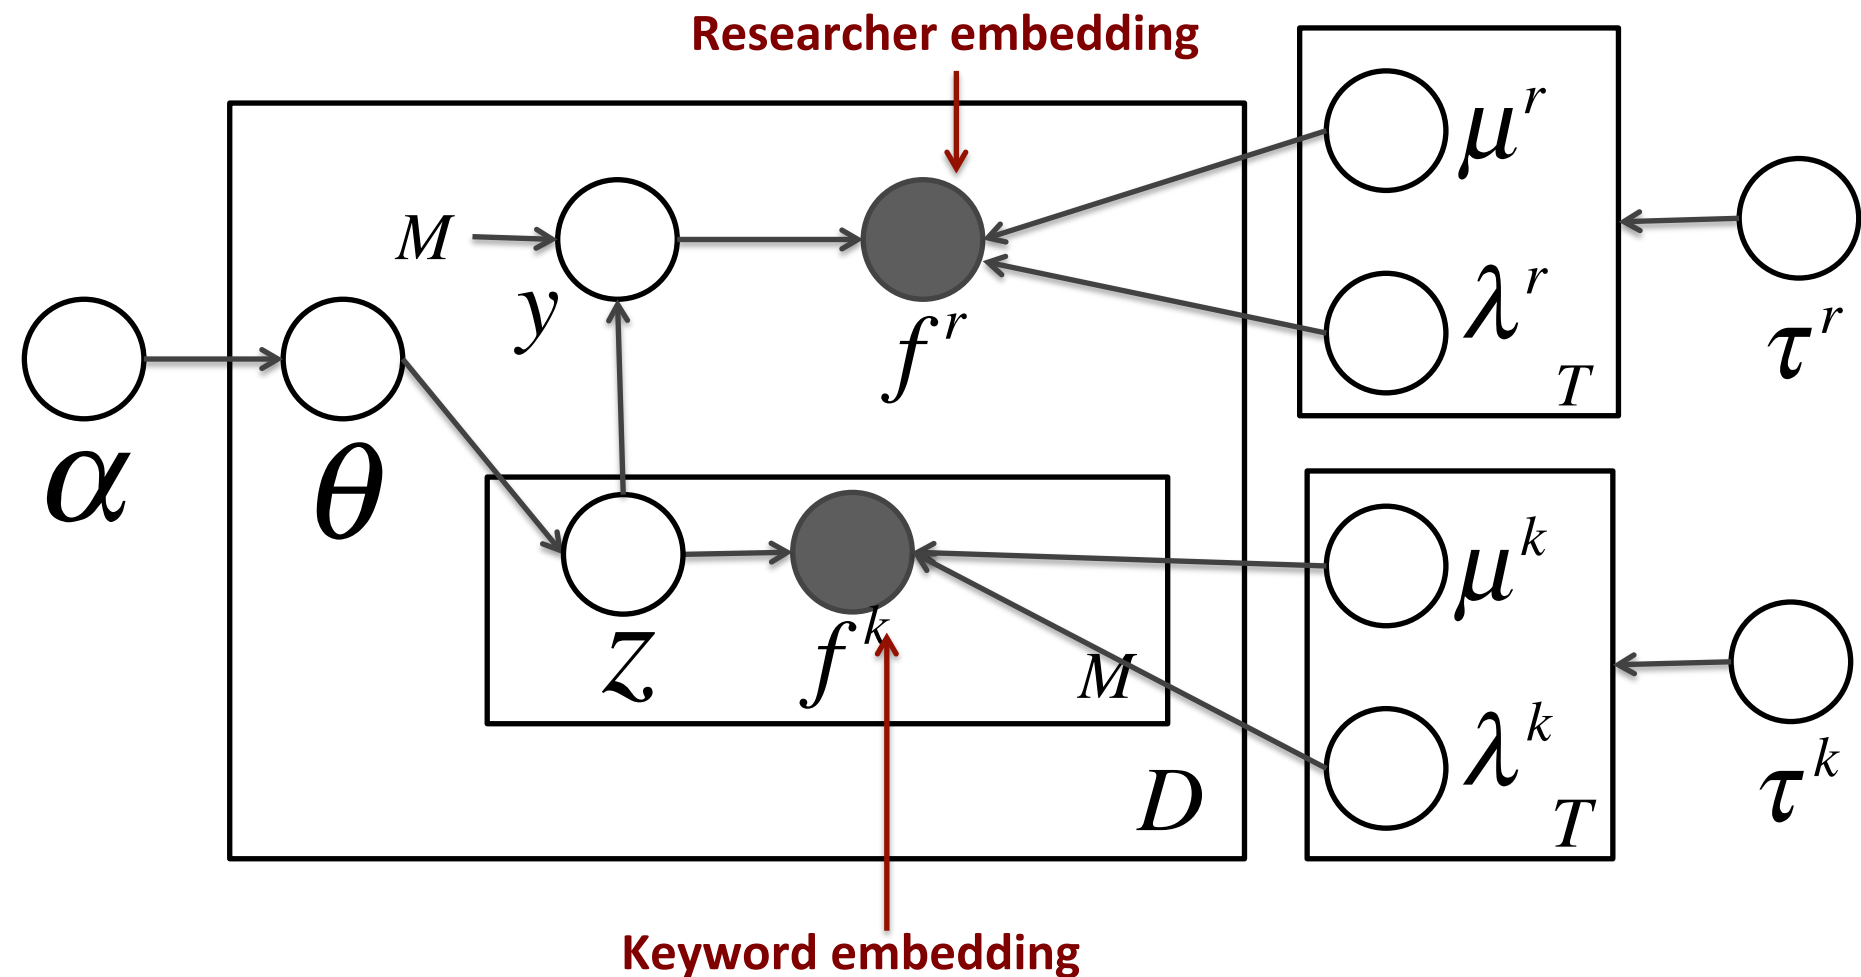


# Multi-source Bayesian embeddings




Number of documents:  $D$ , number of topics:  $T$ , dimension of embedding:  $E$

# Multi-source Bayesian embeddings




Number of documents:  $D$ , number of topics:  $T$ , dimension of embedding:  $E$

# Multi-source Bayesian embeddings



Number of documents:  $D$ , number of topics:  $T$ , dimension of embedding:  $E$


# Multi-source Bayesian embeddings



**Generate Gaussian distribution  
for each topic**

Number of documents:  $D$ , number of topics:  $T$ , dimension of embedding:  $E$


# Multi-source Bayesian embeddings



**Generate the topic for each word**

Number of documents:  $D$ , number of topics:  $T$ , dimension of embedding:  $E$

# Multi-source Bayesian embeddings




**Generate the topic for each user**

Number of documents:  $D$ , number of topics:  $T$ , dimension of embedding:  $E$



# Multi-source Bayesian embeddings



**Generate embeddings for keywords and users**

Number of documents:  $D$ , number of topics:  $T$ , dimension of embedding:  $E$

# Inference

- Collapsed Gibbs sampling
- The joint probability

$$p(\theta, \mu^r, \lambda^r, \mu^k, \lambda^k, z, y, f^r, f^k; \alpha, \tau^r, \tau^k) = p(\theta; \alpha) p(\mu^r, \lambda^r; \tau^r) p(\mu^k, \lambda^k; \tau^k) \\ p(z|\theta) p(f^k|z, \mu^k, \lambda^k) p(f^r|y, \mu^r, \lambda^r) p(y|z)$$

# Inference

Dirichlet distribution

$$p(\theta_d; \alpha) = \frac{1}{\Delta(\alpha)} \prod_{t=1}^T \theta_{dt}^{\alpha_t - 1}$$

Normal Gamma distribution

$$p(\mu_{te}^r, \lambda_{te}^r; \tau_e^r = \{\mu_0, \lambda_0, \alpha_0, \beta_0\}) = \frac{\beta_0^{\alpha_0} \sqrt{\lambda_0}}{\Gamma(\alpha_0) \sqrt{2\pi}} \lambda_{te}^{r \alpha_0 - 1/2} e^{\beta_0 \lambda_{te}^r} e^{-\frac{\lambda_0 \lambda_{te}^r (\mu_{te}^r - \mu_0)^2}{2}}$$

$$p(\mu_{te}^k, \lambda_{te}^k; \tau_e^k = \{\mu_0, \lambda_0, \alpha_0, \beta_0\}) = \frac{\beta_0^{\alpha_0} \sqrt{\lambda_0}}{\Gamma(\alpha_0) \sqrt{2\pi}} \lambda_{te}^{k \alpha_0 - 1/2} e^{\beta_0 \lambda_{te}^k} e^{-\frac{\lambda_0 \lambda_{te}^k (\mu_{te}^k - \mu_0)^2}{2}}$$

# Inference

Generating topics

$$p(z_{dm}|\theta_d) = \theta_{dz_{dm}}$$

$$p(y_d|z_d) = \frac{\sum_{m=1}^{M_d} \mathbb{I}(z_{dm} = y_d) + l}{M_d + Tl}$$

# Inference

Generating embeddings

$$p(f_{dm}^k | z_{dm}, \mu^k, \lambda^k) = \frac{1}{\sqrt{2\pi}} \sqrt{\lambda^k} e^{-\frac{\lambda^k}{2} (f_{dm}^k - \mu^k)^2}$$

$$p(f_{dm}^r | z_{dm}, \mu^r, \lambda^r) = \frac{1}{\sqrt{2\pi}} \sqrt{\lambda^r} e^{-\frac{\lambda^r}{2} (f_{dm}^r - \mu^r)^2}$$

# Inference

## Full Conditional

$$p(y_d = t | y_{-d}, z, f^r, f^k) \propto (n_d^t + l) \prod_{e=1}^{E^r} G'(f^r, y, t, e, \tau^r, d)$$

$$p(z_{dm} = t | z_{-dm}, y, f^r, f^k) \propto (n_d^{y_d} + l)(n_d^t + \alpha_t) \prod_{e=1}^{E^k} G'(f^k, z, t, e, \tau^k, dm)$$

$$G'(f, y, t, e, \tau, d) = \frac{\Gamma(\alpha_n)}{\Gamma(\alpha_{n'})} \frac{\beta_{n'}^{\alpha_{n'}}}{\beta_n^{\alpha_n}} \left( \frac{\kappa_{n'}}{\kappa_n} \right)^{\frac{1}{2}} \frac{(2\pi)^{-n/2}}{(2\pi)^{-n'/2}}$$

# Parameter update

$$\theta_d^t = \frac{n_d^t + \alpha_t}{\sum_{t=1}^T (n_d^t + \alpha_t)}$$

$$\mu_t^k = \frac{\kappa_0 \mu_0 + n \bar{x}}{\kappa_0 + n}$$

$$\lambda_t^k = \alpha_n \beta_n^{-1} = \frac{\alpha_0 + n/2}{\beta_0 + \frac{1}{2} \sum_i (x_i - \bar{x})^2 + \frac{\kappa_0 n (\bar{x} - \mu_0)^2}{2(\kappa_0 + n)}}$$

# Embedding update

$$\frac{\partial L}{\partial f_{de}^r} = \sum_{t=1}^T -\lambda_{te}^r (f_{de}^r - \mu_{te}^r)$$

$$\frac{\partial L}{\partial f_{we}^k} = \sum_{t=1}^T n_w^t (-\lambda_{te}^k) (f_{we}^k - \mu_{te}^k)$$



# Learning framework

- Initialize
- Burn-in
  - Sample topics
- Sampling
  - Sample topics
  - Update parameters
  - Update embeddings

# Experiments

- Comparison methods
  - GenVector: our method
  - GenVector-E: without embeddings
  - GenVector-M: without the model
  - GenVector-R: use weakly-supervision score only
  - AM-base: AMiner previous method
  - CountKG: sort by counts after KG matching
  - Author-topic: Author-topic model
  - NTN: Neural tensor network

# Experiments: homepage matching

Methods	Precision@5
<b>GenVector</b>	<b>77.9402%</b>
GenVector-E	77.8548%
GenVector-M	65.5608%
GenVector-R	72.8549%
AM-base	73.8189%
CountKB	54.4832%
Author-topic	74.4397%
NTN	65.8911%

# Experiments: LinkedIn skill matching

Methods	Precision@5
<b>GenVector</b>	<b>26.8468%</b>
GenVector-E	26.5765%
GenVector-M	24.6695%
GenVector-R	26.3063%
AM-base	24.5195%
CountKB	25.4954%
Author-topic	26.4864%
NTN	24.3243%

# Experiments: human labeling bad cases

Methods	Precision@5
<b>GenVector</b>	<b>98.8%</b>
<b>GenVector-R</b>	<b>99.6%</b>
AM-base	81.2%
Author-topic	98.4%
NTN	92.8%

# Online deployment



Upload

Jiawei Han (韩家炜) ✓

Follow

Department of Computer Science, University of Illinois at Urbana-Champaign

Professor

(217) 333-6903

hanj@cs.uiuc.edu

<http://www.cs.uiuc.edu/~hanj/>


External Links

Update



Research Interests

Data Mining Information Extraction Data Analysis  
Machine Learning Text Mining



1985

1990

1995

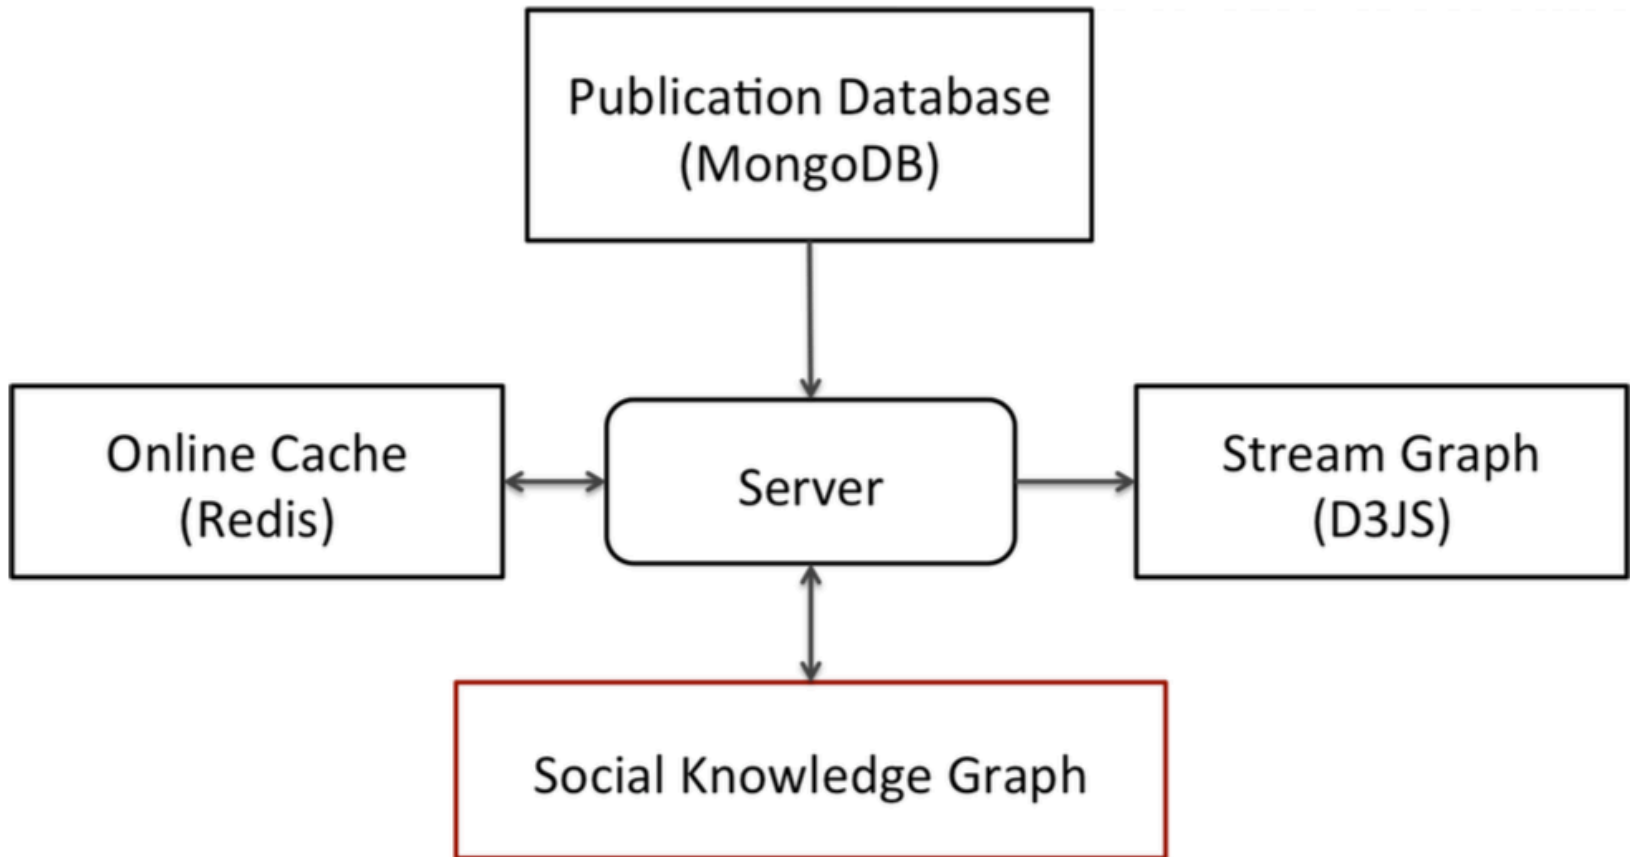
2000

2005

2010

2014

# Online deployment

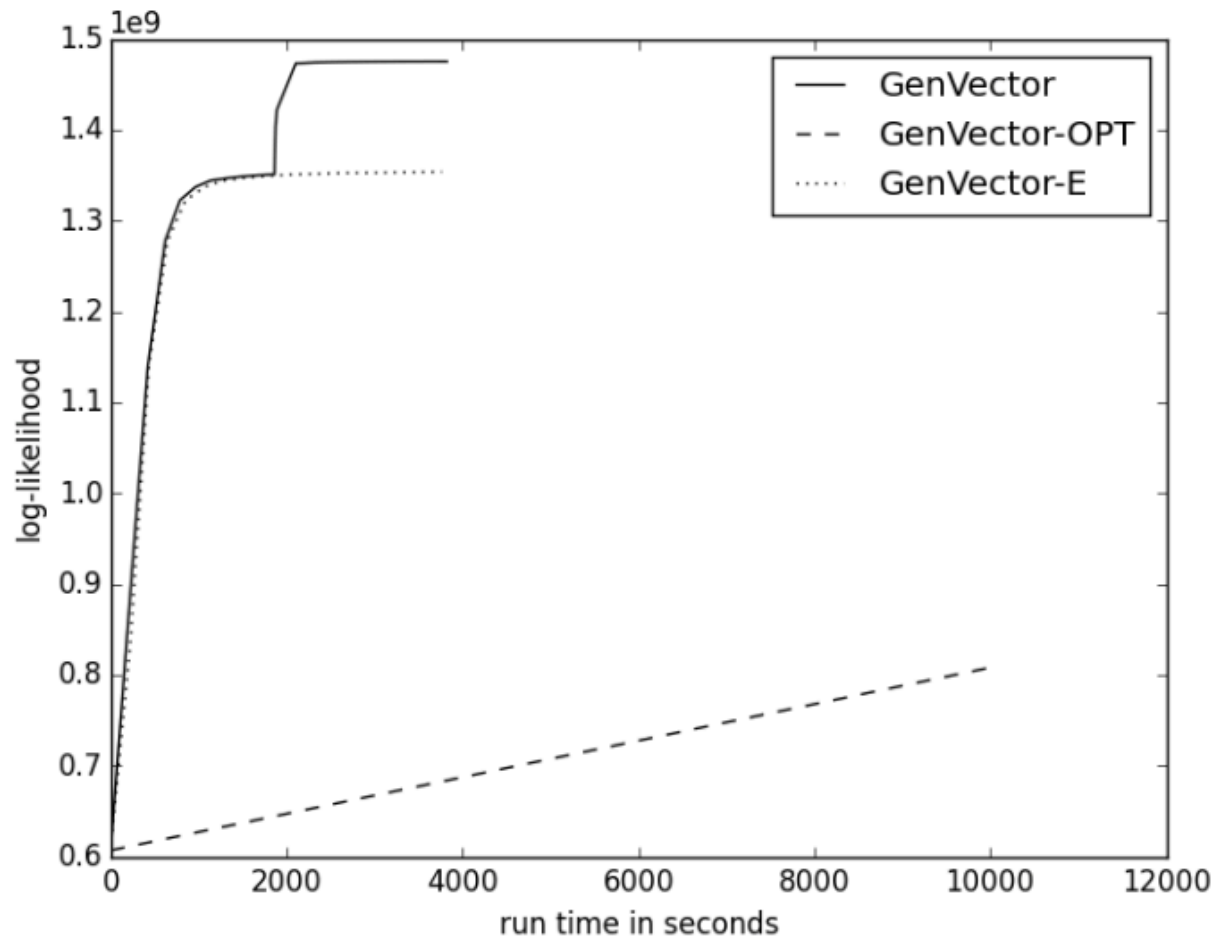


# Implementation optimization

- Faster computation of  $G'()$
- Faster computation of log, exp and pow
- Local variables instead of in-array access
- Multi-thread parallelization



# Run time and convergence



# Online AB-test

✕

Does Jiawei Han have these skills or expertise?

Text Mining ✕

Efficient Algorithm ✕

Data Mining ✕

Machine Learning ✕

Large Databases ✕

Knowledge Discovery ✕

Type another area of expertise

Skip

Vote

Leverage collective intelligence

- evaluate the methods
- leverage user feedback to improve the model

# Online AB-test

Methods	Precision@10
GenVector	96.67%
AM-base	90.00%

# Case study: Andrew Ng

GenVector	AM-base
Unsupervised learning	Challenging problem
Feature learning	Reinforcement learning
Bayesian networks	Autonomous helicopter
Reinforcement learning	Autonomous helicopter flight
Dimensionality reduction	Near-optimal planning

# Case study: Dan Klein

GenVector	AM-base
Language models	Machine translation
Markov models	Word alignment
Probabilistic models	Bleu score
Natural language	Best result
Coreference resolution	Language model

# Case study: Xiaoou Tang

GenVector	AM-base
Feature extraction	Face recognition
Image segmentation	Face image
Image matching	Novel approach
Image classification	Line drawing
Face recognition	Discriminant analysis

# Take-away


- Large-scale
  - Link **38,049,189** researchers to **35,415,011** knowledge concepts
- Fast
  - **60** times speed up
- Accurate
  - Decrease the error rate by **67%** online
- Novel
  - Bridge social networks and collective knowledge
  - bridge topic models and network/word embedding
- Real-world impact
  - Online service with **183,876** visits

# Appendix



# Learning keyword embeddings

- Skip-gram



# Learning keyword embeddings

- Skip-gram
  - Use the current keyword to predict the context
  - Objective function

$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t)$$

$$p(w_O | w_I) = \frac{\exp \left( v'_{w_O}{}^\top v_{w_I} \right)}{\sum_{w=1}^W \exp \left( v'_w{}^\top v_{w_I} \right)}$$

# Learning keyword embeddings

- Scan through all titles and abstracts
  - Extract n-grams according to Wikipedia concepts
- Replace all extracted n-grams in the Wikipedia corpus as a token
  - E.g., machine learning -> machine\_learning
- Train a skip-gram model on the processed corpus

# Learning network embeddings

- DeepWalk
  - Generate a random walk sequence from each node
  - Train a skip-gram model on the random walk sequence

# Weakly supervision

- Given a researcher, extract all the keywords in his papers' titles, denoted as  $k_1, k_2, \dots, k_n$ .
- Let  $c_i$  be the count of the keyword  $k_i$  in the author's papers' titles.
- Compute a score for each keyword  $k_i$

$$s_i = \sum_j c_j \cos_{i,j}$$

- Select top-k keywords as weakly-supervised information